

Global Volatility Spillovers with the BVAR Package

Tomas Samaj Luuk Vos

15 April 2026

WU Vienna

Main Results

Project Overview

What is our project?

We study how volatility shocks propagate across six major equity markets using a Bayesian VAR, and we compare calm and crisis regimes to see whether global transmission intensifies under stress.

Empirical design

- Six implied-volatility indices in levels
- Full sample plus calm-vs-crisis comparison
- Generalized IRFs, GFEVDs, and conditional forecasts as the main outputs

Why this is a good BVAR application

- Volatility is persistent and strongly comoving.
- A six-variable VAR with five lags already has 186 coefficients.
- Bayesian shrinkage is therefore central to making the model usable.

Our angle

We use this project to show how the BVAR package turns a theoretically rich model into an applied workflow that is easy to estimate, diagnose, and interpret.

Why BVAR?

Dimensionality problem

A VAR with M variables and p lags has

$$M + M^2 p$$

slope parameters before the covariance matrix.

With $M = 6$ and $p = 5$, this already implies

$$6 + 6^2 \cdot 5 = 186$$

coefficients to regularise.

Why Bayesian shrinkage helps

- It shrinks weak coefficients toward sensible benchmarks.
- It reduces overfitting in a medium-size system.
- It improves stability for forecasting and impulse analysis.

Practical message

For this kind of volatility system, Bayesian regularisation is not just elegant theory; it is what makes the VAR empirically tractable.

Compact Theory: VAR and Minnesota Prior

Bayesian VAR

$$\mathbf{y}_t = \mathbf{a}_0 + A_1 \mathbf{y}_{t-1} + \cdots + A_p \mathbf{y}_{t-p} + \boldsymbol{\varepsilon}_t, \quad \boldsymbol{\varepsilon}_t \sim \mathcal{N}(\mathbf{0}, \Sigma)$$

Minnesota prior

$$\mathbb{E}[(A_s)_{ij} \mid \Sigma] = \begin{cases} 1 & i = j, s = 1 \\ 0 & \text{otherwise} \end{cases}$$

$$\text{Var}[(A_s)_{ij} \mid \Sigma] \propto \frac{\lambda^2 \psi_j}{s^{2\alpha}} \cdot \frac{\sigma_{ii}}{\sigma_{jj}}$$

Interpretation

- Each series is assumed to be persistent in its own first lag.
- Cross-variable lag effects are shrunk toward zero unless the data strongly support them.
- λ controls overall tightness, α controls lag decay, and ψ_j scales variable-specific uncertainty.

Why this fits our data

Volatility indices are persistent and correlated, so the Minnesota prior stabilises the system without ruling out spillovers.

Compact Theory: Learning Shrinkage from the Data

Hierarchical prior selection

$$p(\theta, \gamma | \mathbf{y}) \propto p(\mathbf{y} | \gamma) p(\gamma), \quad p(\mathbf{y} | \gamma) = \int p(\mathbf{y} | \theta, \gamma) p(\theta | \gamma) d\theta$$


Interpretation

- Hyperparameters such as λ are not fixed manually.
- They are learned through the marginal likelihood.
- The procedure automatically balances fit against overfitting.

Why BVAR helps

- The package estimates the model, samples hyperparameters, and produces diagnostics in one workflow.
- Generalized IRFs, GFEVDs, and scenario forecasts all come from the fitted object and its posterior draws.
- This makes the methodology much easier to operationalise in practice.

Data and Empirical Setting



Setup

- VIX, V2X, IVIUK, VNKY, VHSI, INVIXN
- Sample: 2015-01-05 to 2026-03-31
- Calm sample: 2015-01-05 to 2019-12-31
- Crisis sample: 2020-03-02 to 2020-12-31

What the figure shows

- All markets spike sharply in early 2020.
- Levels differ by region, but common movement is strong.
- This makes the system well suited to a spillover analysis.

BVAR Package Workflow

```
mn <- bv_minnesota(  
  lambda = bv_lambda(mode = 0.2, sd = 0.4),  
  alpha = bv_alpha(mode = 2),  
  psi = bv_psi(scale = 0.004, shape = 0.004)  
)  
  
fit <- bvar(  
  data = data_full, lags = 5,  
  n_draw = 25000, n_burn = 10000,  
  priors = bv_priors(hyper = c("lambda", "alpha", "psi"), mn = mn),  
  irf = bv_irf(horizon = 20, fevd = TRUE, identification = FALSE),  
  fcast = bv_fcast(horizon = 20)  
)  
  
irf_full <- irf(fit, horizon = 20, conf_bands = c(0.05, 0.16))  
gfevd_full <- compute_gfevd(irf_full, fit$sigma, var_names, horizon = 20)  
pred_full <- predict(fit, conf_bands = c(0.05, 0.16))
```


What this workflow gives us

- One object for estimation, diagnostics, and post-estimation
- Generalized IRFs and GFEVD-ready posterior draws
- Conditional forecasts for stress scenarios

Why this matters

The package makes the full Bayesian workflow transparent enough for teaching and gives us the objects needed for ordering-invariant post-estimation analysis.

Minnesota-Only Convergence Diagnostics




Takeaway

- λ and α stabilise in clear regions.
- Some variable-specific scales mix more slowly, which is normal in a six-variable hierarchical system.
- Overall, the baseline Minnesota specification looks usable.

Full-Sample Generalized IRFs: 1 SD VIX Shock

IRF: Response to a 1 SD VIX Shock - Full Sample


Sample: 2015-01-05 to 2026-03-31 | 68% and 90% posterior credible bands | Cholesky identification



What we learn

- Without imposing a recursive ordering, every market still reacts positively to a VIX shock.
- Europe, the UK, and Japan respond strongly on impact.
- India reacts more gradually, but the response remains positive over the full horizon.
- The slow decay is consistent with persistent volatility dynamics.

Full-Sample GFEVD



What we learn

- The ordering-invariant GFEVD still points to VIX as the main external shock in most panels.
- V2X and IVIUK remain especially exposed to US volatility.
- Japan, Hong Kong, and India keep more mixed local and regional contributions.
- GFEVD complements the generalized IRFs by showing which shocks matter quantitatively without imposing causal ordering.


Crisis vs Calm: Generalized IRFs

IRF Comparison Across Regimes

Calm sample: 2010-01-05 to 2019-12-31 | Crisis sample: 2020-03-02 to 2020-12-31


IRF: Response to a 1 SD VIX Shock - Calm Period

Sample: 2010-01-05 to 2019-12-31 | 68% and 95% posterior credible bands



IRF: Response to a 1 SD VIX Shock - COVID Acute Crisis


Sample: 2020-03-02 to 2020-12-31 | 68% and 95% posterior credible bands



What changes in crisis?

- Crisis responses are larger on impact and more uncertain.
- Europe and the UK react much more strongly during 2020-03-02 to 2020-12-31 than in 2015-01-05 to 2019-12-31.
- India becomes far more responsive and peaks later, suggesting delayed but substantial transmission.
- The positive sign survives without imposing a recursive ordering, but the magnitude depends strongly on market conditions.

Crisis vs Calm: GFEVD




What changes in crisis?

- The crisis period shifts Hong Kong and India away from own or regional shares toward VIX, V2X, and VIUK.
- Japan also becomes less self-driven as global shocks take more of the variance.
- VIX remains important, but the crisis decomposition becomes more broadly global rather than locally concentrated.

Interpretation

Ordering-invariant variance shares show broader contagion in crisis periods: local components shrink and common external stress explains more of the system.

Conditional Forecast: Persistent VIX Stress




What we learn

- Fixing VIX one forecast standard deviation above baseline lifts every other forecast path.
- V2X, IVIUK, and VNKY show the largest upward shifts.
- The gap to baseline persists over the full horizon, so this is not just a one-day response.

Why this showcases the package

The scenario is implemented directly with `bv_fcast()` and `predict()`, so stress testing becomes part of the standard workflow.

Conditional Forecast: India Stress Scenario




What we learn

- A one-standard-deviation India shock still raises the other paths, but much less than the VIX scenario.
- This is interesting because India is the only emerging-market index in the system.
- The figure suggests weaker outward transmission from India than from the US.

Interpretation

India is affected by global volatility, but its own stress scenario appears to be less system-defining than a comparable VIX shock.

Robustness to Lag Length



What we learn

- The ranking of responses is stable across 1, 3, and 5 lags.
- Shorter lag models smooth the early dynamics.
- Five lags allow sharper short-run spillovers, but the economic message stays the same.

Bottom line

Our conclusions are not an artefact of one particular lag choice.

SOC and SUR: What Do They Add?

Sum-of-coefficients prior (SOC)

$$\mathbf{y}^+ = \text{diag}\left(\frac{\bar{\mathbf{y}}}{\mu}\right), \quad \mathbf{x}^+ = [0, \mathbf{y}^+, \dots, \mathbf{y}^+]$$


Single-unit-root prior (SUR)

$$\mathbf{y}^{++} = \frac{\bar{\mathbf{y}}^\top}{\delta}, \quad \mathbf{x}^{++} = \left[\frac{1}{\delta}, \mathbf{y}^{++}, \dots, \mathbf{y}^{++}\right]$$

Interpretation

- SOC pushes the system toward persistent no-change dynamics.
- SUR allows common stochastic trends or cointegration-like persistence.
- Because our series are volatility levels, these priors are admissible as a robustness extension.

Minnesota + SOC + SUR: Convergence Diagnostics




Takeaway

- The marginal likelihood, λ , SOC, and SUR fluctuate around stable regions.
- Posterior densities are reasonably concentrated.
- The extended specification also converges in a usable way.

SOC/SUR Traces and Why the Results Barely Change

Hyperparameter Trace and Density Diagnostics - Full Sample

Sample: 2010-01-01 to 2020-03-01 | Posterior prior: Minnesota + SOC + SUR | Geweke z outside +/-1.96 suggests non-stationarity



What the traces show

- The added hyperparameters stay in stable regions.
- Posterior mass concentrates around moderate values for λ , SOC, and SUR.

Why the economic results barely change

- The data are already persistent and estimated in levels.
- The Minnesota prior already captures most of the relevant persistence.
- SOC and SUR mainly reinforce that persistence rather than introduce a new cross-market story.

Project Conclusions

- A US volatility shock propagates positively to all markets in the system, so spillovers are clearly global.
- The transmission is stronger and more uncertain during the 2020 crisis window than in the calm period.
- The GFEVD still highlights VIX as the key external driver, while Hong Kong and India retain relatively larger local or regional components.
- Conditional forecasting confirms this asymmetry: a VIX stress scenario moves the whole system more than an India stress scenario.
- Our main findings are robust to lag choice and remain qualitatively unchanged when SOC and SUR are added.

What This Shows About the BVAR Package

- BVAR turns a full Bayesian workflow into a compact set of functions: prior setup, estimation, diagnostics, generalized IRFs, GFEVD-ready posterior objects, and scenario forecasts.
- The package makes hierarchical shrinkage operational rather than purely theoretical, which is especially valuable in medium-size VAR systems.
- It is strong not only for estimation, but also for communication: the same fitted object supports generalized responses, ordering-invariant variance shares, robustness checks, and stress scenarios.
- Our project is a good example of this strength: the package lets us move from model setup to empirical conclusions in a coherent and reproducible way.

Appendix: VAR Setup

Model

$$\mathbf{y}_t = \mathbf{a}_0 + A_1 \mathbf{y}_{t-1} + \cdots + A_p \mathbf{y}_{t-p} + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(\mathbf{0}, \Sigma)$$

Why OLS is not enough

- Parameter count rises quickly with the number of variables and lags.
- Financial volatility series are persistent and strongly correlated.
- Unregularised VARs can overfit in-sample noise.

Why Bayesian shrinkage helps

- Priors shrink weak coefficients toward economically sensible benchmarks.
- The data can still dominate when the signal is strong.
- This improves stability and forecast performance in medium-size systems.

Appendix: Minnesota Prior

Prior mean

$$\mathbb{E}[(A_s)_{ij} \mid \Sigma] = \begin{cases} 1 & \text{if } i = j, s = 1 \\ 0 & \text{otherwise} \end{cases}$$

Prior variance

$$\text{Var}[(A_s)_{ij} \mid \Sigma] \propto \frac{\lambda^2 \psi_j}{s^{2\alpha}} \cdot \frac{\sigma_{ii}}{\sigma_{jj}}$$

- λ : overall tightness
- α : lag decay
- ψ_j : variable-specific scaling
- Intuition: own first lags are persistent; cross-lag effects are shrunk strongly toward zero

Appendix: Hierarchical Hyperparameter Selection

Key idea

Treat the prior hyperparameters as random variables:

$$p(\boldsymbol{\theta}, \boldsymbol{\gamma} \mid \mathbf{y}) \propto p(\mathbf{y} \mid \boldsymbol{\gamma}) p(\boldsymbol{\gamma}), \quad p(\mathbf{y} \mid \boldsymbol{\gamma}) = \int p(\mathbf{y} \mid \boldsymbol{\theta}, \boldsymbol{\gamma}) p(\boldsymbol{\theta} \mid \boldsymbol{\gamma}) d\boldsymbol{\theta}$$

Why this matters

- Shrinkage is learned from the data instead of fixed by hand.
- The marginal likelihood penalises overly flexible models automatically.
- Results become easier to justify and replicate.

How BVAR implements it

- Conjugate NIW structure gives closed-form posteriors for core parameters.
- Hyperparameters are sampled with Metropolis-Hastings.
- Diagnostics are available directly from the fitted object.

Appendix: SOC and SUR Dummy Priors

Sum-of-coefficients prior (SOC)

$$\mathbf{y}^+ = \text{diag}\left(\frac{\bar{\mathbf{y}}}{\mu}\right), \quad \mathbf{x}^+ = [0, \mathbf{y}^+, \dots, \mathbf{y}^+]$$

Single-unit-root prior (SUR)

$$\mathbf{y}^{++} = \frac{\bar{\mathbf{y}}^\top}{\delta}, \quad \mathbf{x}^{++} = \left[\frac{1}{\delta}, \mathbf{y}^{++}, \dots, \mathbf{y}^{++}\right]$$

- SOC shrinks the system toward “no large changes from current levels.”
- SUR allows persistent common movement and is useful when cointegration-like behaviour is plausible.
- These priors are designed for VARs in levels, which is why we only use them as a level-data robustness extension.

Appendix: Estimation Choices in Our Script

Baseline specification

- Six volatility indices in levels
- Lags: $p = 5$
- Draws: 25,000
- Burn-in: 10,000
- IRF / forecast horizon: 20 trading days

Identification and samples

- No recursive ordering is imposed.
- `bv_irf(..., identification = FALSE)`
- GFEVD follows Pesaran-Shin and is ordering-invariant.
- Calm sample: 2015-01-05 to 2019-12-31
- Crisis sample: 2020-03-02 to 2020-12-31

Interpretation note

All results should be read as reduced-form, ordering-invariant volatility transmission patterns rather than structural causal effects.

References i

-  Kuschnig, N., & Vashold, L. (2021).
BVAR: Bayesian Vector Autoregressions with Hierarchical Prior Selection in R.
Journal of Statistical Software, 100(14), 1–27.
-  Giannone, D., Lenza, M., & Primiceri, G. E. (2015).
Prior Selection for Vector Autoregressions.
The Review of Economics and Statistics, 97(2), 436–451.
-  Litterman, R. B. (1980).
Techniques of forecasting using vector autoregressions.
Federal Reserve Bank of Minneapolis Working Paper 115.
-  Doan, T., Litterman, R., & Sims, C. (1984).
Forecasting and conditional projection using realistic prior distributions.
Econometric Reviews, 3(1), 1–100.



Sims, C. A. (1993).

A nine-variable probabilistic macroeconomic forecasting model.

In Business Cycles, Indicators and Forecasting.



Sims, C. A., & Zha, T. (1998).

Bayesian methods for dynamic multivariate models.

International Economic Review, 39(4), 949–968.



Pesaran, M. H., & Shin, Y. (1998).

Generalized impulse response analysis in linear multivariate models.

Economics Letters, 58(1), 17–29.